Aeolus Renewables

Drivetrain condition assessment using 10-minute SCADA measurements

131,760
sensor records

15
turbines

3.00%
recorded faults

Predictive Modeling with Machine Learning and Neural Networks | Low-Code Path

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Contents / Agenda

01 Executive summary and business problem

02 Exploratory analysis and preprocessing

03 Model comparison and final selection

04 Actionable recommendations

A Appendix: data, baselines, tuning and evidence

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Executive Summary

VALIDATION RECALL	TEST RECALL	TEST PRECISION	TEST F1
92.49%	**86.40%**	**49.48%**	**62.93%**
Tuned ANN selected	1,245 / 1,441 faults	False-alert tradeoff	Balance of P and R

The operational decision

Use the ANN as analyst decision support, not an automatic turbine shutdown trigger.

1,271 false-positive records require remote triage; 196 fault records were missed.

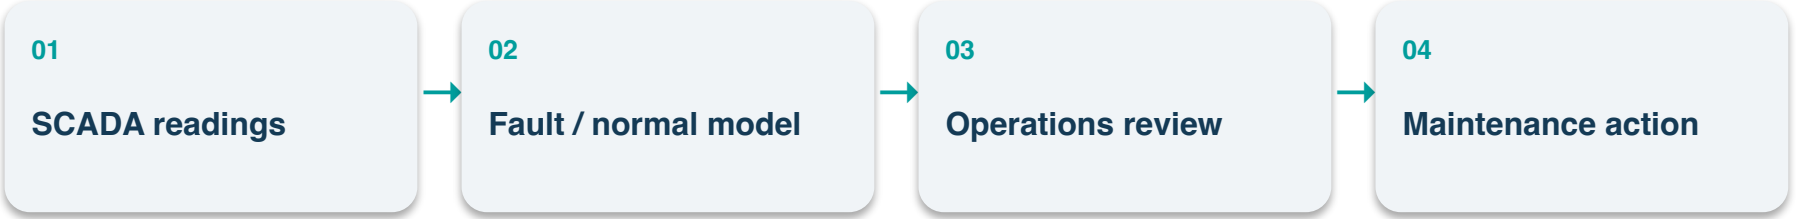
Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Business Problem Overview and Solution Approach

WHY IT MATTERS

Missed drivetrain faults can lead to extended downtime, emergency repairs and lost energy production.

CONDITION-BASED WORKFLOW

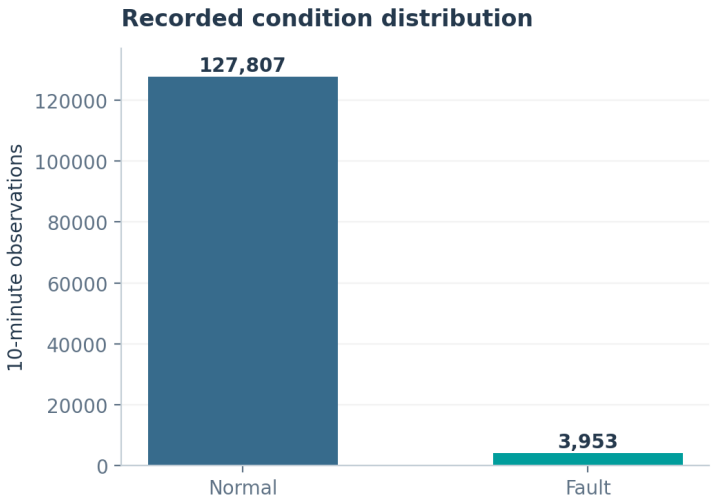


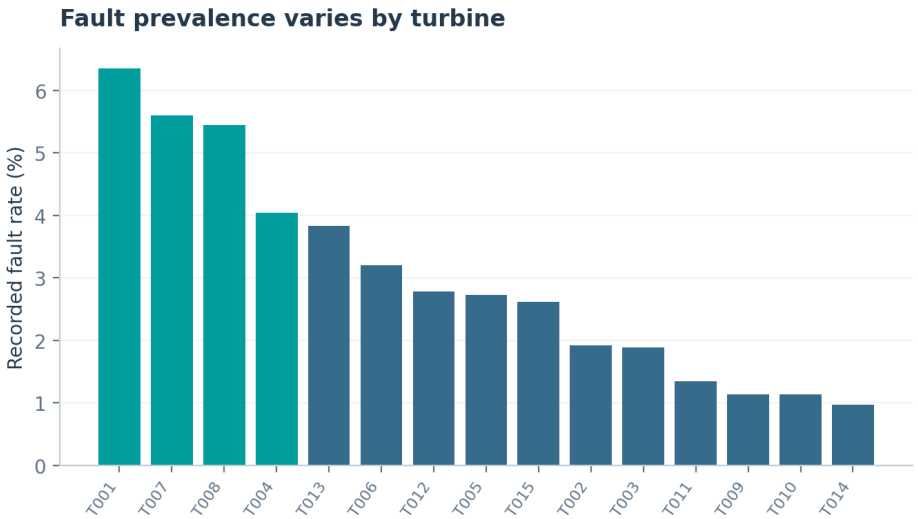
Method: timestamp-grouped chronological split, training-only preprocessing, five baseline models, two tuned candidates, validation-driven selection.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

EDA Results — Class Balance and Asset Patterns

131,760 readings • 15 turbines • 3.00% recorded faults

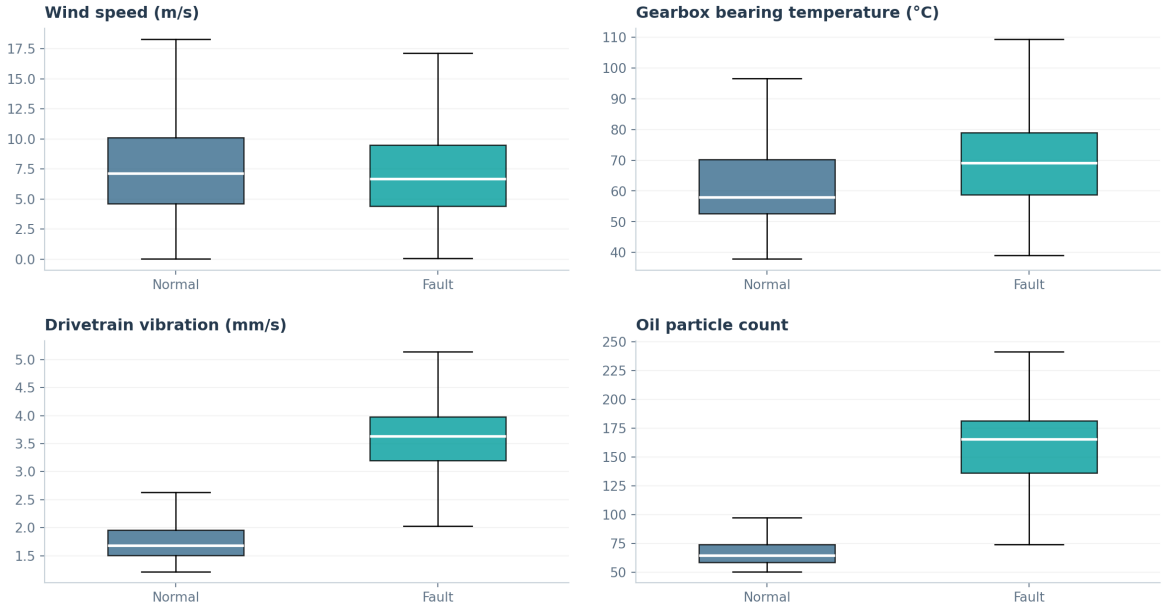




The minority fault class makes accuracy alone misleading; turbine-level rates support prioritised investigation.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

EDA Results — Sensor Evidence



KEY OBSERVATIONS

Vibration, gearbox temperature and oil particles are visibly elevated in fault-state records.

Wind speed alone provides weaker separation, supporting multi-signal evaluation.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Data Preprocessing and Leakage Prevention

DATA QUALITY

0 duplicate rows; missing gearbox-oil (645), generator-bearing (690), oil-pressure (636).

TREATMENT

Training-median imputation; ANN StandardScaler fit using training rows only.

FEATURE ENGINEERING

Time cycles, power utilisation and temperature-rise variables.

EVALUATION DESIGN

Timestamp-grouped chronological 70% / 15% / 15%; test reserved until final evaluation.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Model Performance Summary and Final Selection

Tuned model	Validation recall	F1	Precision
ANN (selected)	92.49%	53.54%	37.67%
XGBoost	69.10%	69.34%	69.57%

Selection rule

Maximise validation recall; F1 and precision are secondary tie-breakers.

ANN catches more recorded fault states. XGBoost offers better precision and F1, so the false-alert tradeoff must be managed.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Actionable Business Recommendations

01 Triage

Review persistent model alerts and corroborating sensor signals remotely.

02 Dispatch

Prioritise technician inspections when risk persists across successive intervals.

03 Learn

Link alerts to work orders; track false alarms, missed faults, and operational workload.

04 Validate

Run prospective field validation before promising lead time or cost savings.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

SCADA panel

8,784 timestamp groups, recorded at 10-minute intervals across 15 turbines.

Environment

Wind, ambient conditions

Operation

Power and rotational speeds

Condition

Bearing heat, vibration, oil particles

Asset history

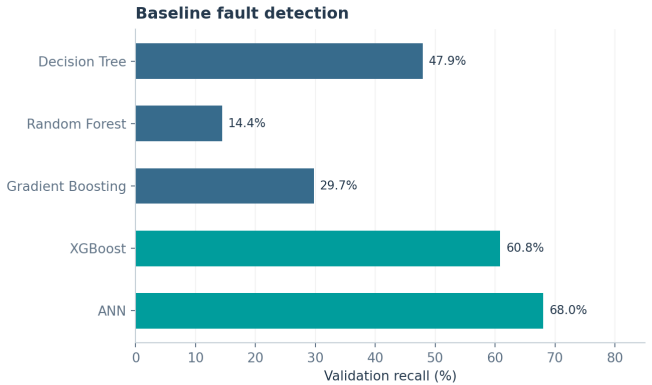
Hours, load cycles and maintenance

Target: failure = 1 denotes a recorded fault condition, not a future failure horizon.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Model Building — Five Baseline Classifiers

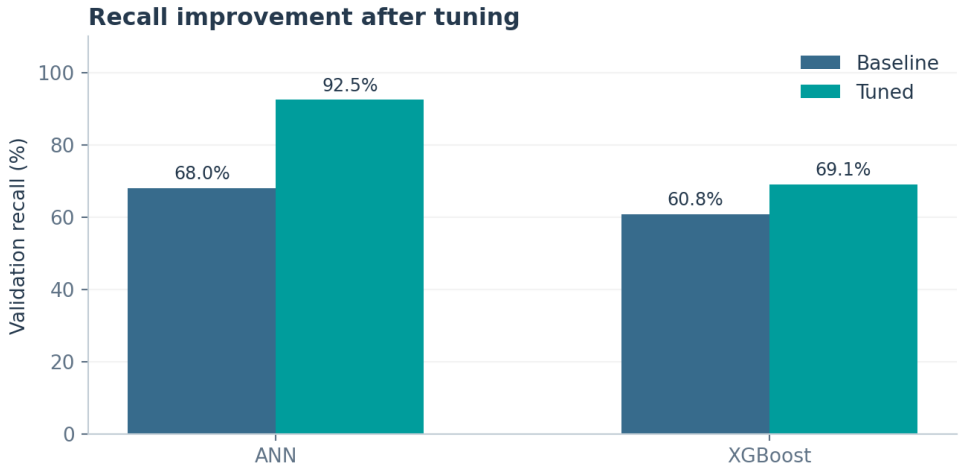
Model	Accuracy	Precision	Recall	F1
Decision Tree	90.70%	27.74%	47.93%	35.14%
Random Forest	95.35%	82.87%	14.44%	24.59%
Gradient Boosting	94.54%	46.96%	29.74%	36.42%
XGBoost	96.64%	71.01%	60.83%	65.53%
ANN	94.79%	50.28%	68.05%	57.83%



ANN and XGBoost have the highest baseline recall and were selected for tuning.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Model Improvement — Hyperparameter Tuning



XGBOOST

6 sampled configurations; 2 timestamp-grouped expanding folds; recall objective.

Best CV recall 48.45%; holdout validation recall 69.10%.

ANN

64/24 hidden units; BatchNorm, dropout 0.30/0.15, Adam, early stopping.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Final Model Performance — Untouched Test

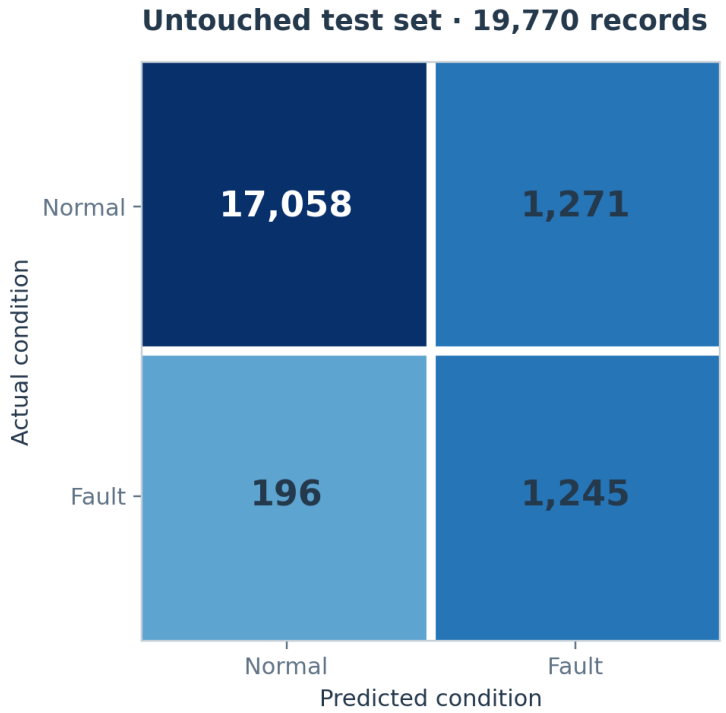
TEST RECALL
86.40%

TEST PRECISION
49.48%

TEST ACCURACY
92.58%

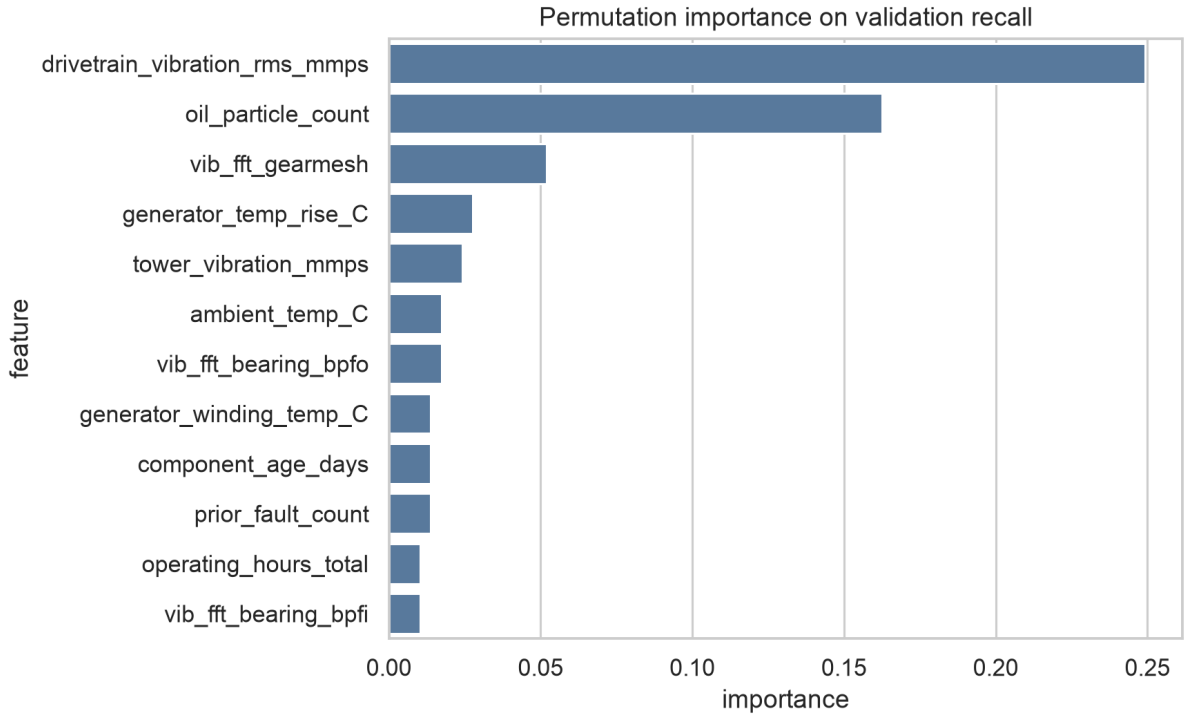
TEST F1
62.93%

1,245 detected • 196 missed
1,271 false alerts • 17,058 true normals



Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Feature Interpretation and Limitations



KEY DRIVERS

Drivetrain vibration and oil particle count show the largest validation-recall impacts.

Permutation importance measures association, not causation or a future-failure lead time.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Timestamp-Grouped Split Verification

TRAIN

92,220

6,148 timestamps

VALIDATION

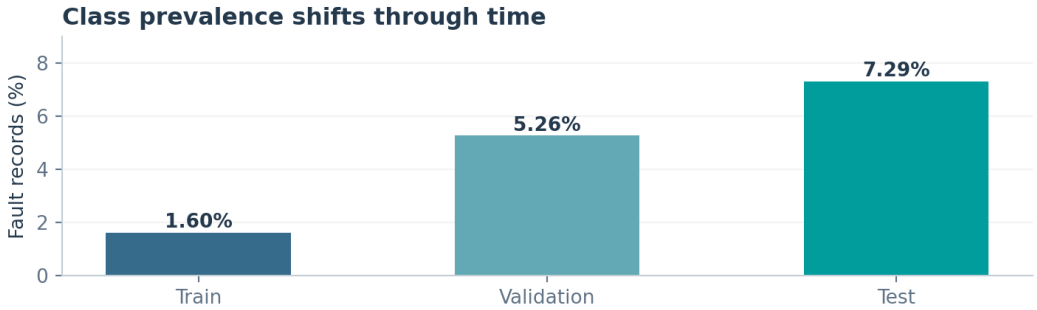
19,770

1,318 timestamps

TEST

19,770

1,318 timestamps



LEAKAGE CHECK

All 15 turbine records at a timestamp remain in the same partition. Failure prevalence changes over time.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

Operational Use Case and Next Steps

STAGE 1

Model flags current abnormal condition.

STAGE 2

Analyst checks persistence and supporting sensors.

STAGE 3

Maintenance team prioritises planned inspection.

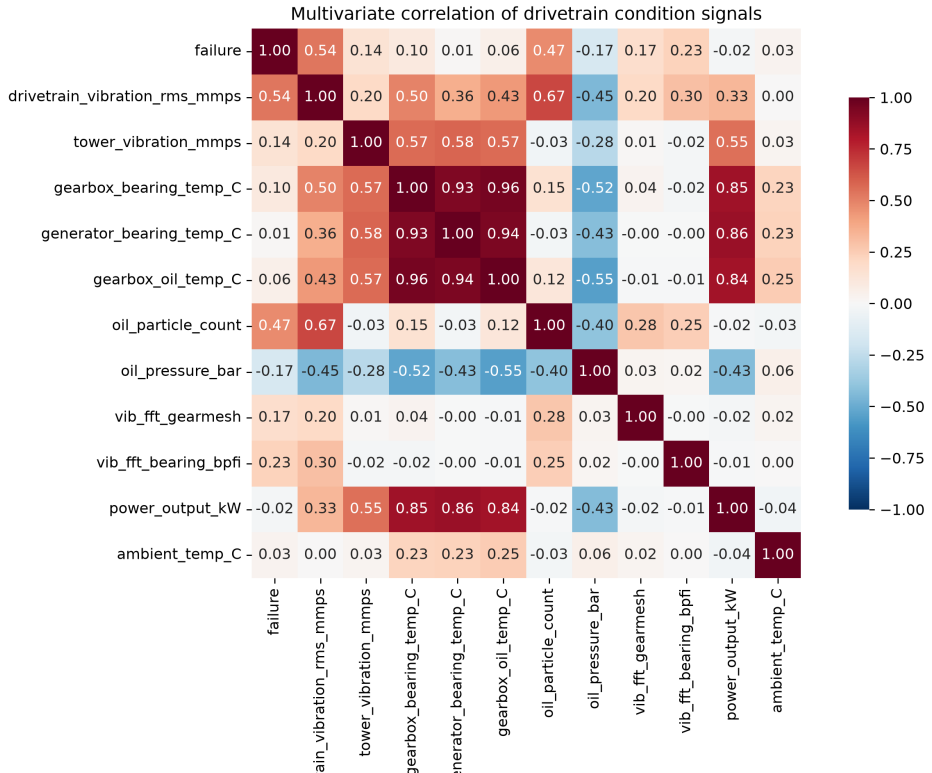
DEPLOYMENT GUARDRAILS

Retrospective proof of concept: prospectively validate alert value, technician workload, model drift and downtime outcomes before automated action or savings claims.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.

EDA Results — Multivariate Sensor Relationships

Correlation heatmap uses actual SCADA values across thermal, vibration, oil and operating variables.



What the data shows

- Drivetrain vibration has the s
correlation with recorded fault
- Oil-particle count is also mat
bearing-frequency vibration is
- Oil pressure is negatively rel
bearing temperature is weake

Business implication

No single sensor separates ev
channel model combines com
and reduces dependence on r
thresholds.

tion prohibited.

Tuning Comparison — Recall First, Precision Trade-

Both candidates were chosen from baseline validation evidence. Selection used holdout valida
results were excluded.

Model	Baseline P / R / F1	Tuned P / R / F1	
ANN	50.28% / 68.05% / 57.83%	37.67% / 92.49% / 53.54%	
XGBoost	71.01% / 60.83% / 65.53%	69.57% / 69.10% / 69.34%	

e

Why ANN is selected
ANN has the highest validation recall (92.49%), so it detects more recorded fault states
business priority. XGBoost has higher precision and F1, making it a credible lower-alert-
This trade-off is documented rather than hidden.

Proprietary content. © Great Learning. All Rights Reserved. Unauthorized use or distribution prohibited.